

Kunj Chetan Mehta

New Brunswick, NJ | +1 8484371589 | kunjmehta@gmail.com | [kunjmehta.github.io](https://github.com/kunjmehta) | linkedin.com/in/kunjmehta

EDUCATION

Master of Science in Computer Science | Rutgers University, New Jersey | GPA – 3.91/4 **Expected May 2023**
Computer Science Exchange Program | Princeton University, New Jersey **Sep 2022 – Dec 2022**
Bachelor of Technology in Computer Engineering | Mumbai University, India | CGPA – 8.52/10 **Oct 2020**
Relevant Coursework – Artificial Intelligence (AI), Machine Learning (ML), Data Warehousing and Mining, Distributed Cloud Computing, Digital Image Processing, Advanced Database Management System, Analysis of Algorithms, Data Structures, Operating System, Blockchain Technology, Mathematical Foundations of Data Science, Advanced Computer Vision

TECHNICAL SKILLS

- **Development:** Python, Java, SQL, PySpark, R, C++, JavaScript, Android, MS Excel
- **Frameworks and Libraries:** NumPy, pandas, matplotlib, seaborn, scikit-learn, MLlib, streamlit, Pytorch, Deep Graph Library, NLTK, Pillow, OpenCV, Flask
- **Cloud Services:** AWS Sagemaker, AWS EMR, AWS Lambda, Google Big Query, Docker, Apache Airflow, Kafka
- **Databases and Visualization:** MySQL, AWS Redshift, Amazon RDS, Tableau, Power BI, Looker
- **Machine Learning:** Regression, Decision Tree, Clustering, EDA, CNN, NLP, Time Series, Ensemble & Multimodal Learning, Recommender Systems
- **Certifications:** Cloud Practitioner (AWS), LookML Developer (Looker), Machine Learning Engineer (Udacity), Deep Learning Specialization (Coursera), Applied Data Science with Python Specialization (Coursera) [\[link\]](#)

EXPERIENCE

- Data Scientist Intern** | Eluvio | Berkeley, CA **Jun 2022 – Aug 2022**
- Part of the machine learning team building the media meta-tagging framework for media distribution on the blockchain
 - Engineered the logo detection and classification pipeline from supervised to zero-shot learning paradigm. Reduced the number of false logo detections by 8 percentage points
 - Pioneered a NFT recommender system end-to-end. Led the creation of a near real-time ETL pipeline to ingest model-ready blockchain data for training. Deployed a test MVP of the recommender handling 3000 concurrent users efficiently
 - Created an external data collection pipeline for the movie speech recognition project eventually leading to 35 percentage point decrease in word error rate
- Teaching Assistant** | Rutgers University | New Brunswick, NJ **Sep 2021 – Present**
- Taught R, SQL and Amazon Redshift and graded weekly assignments and exams for 150 students across two courses – “Data 101” and “Database Systems for Data Science”
- Business Analyst** | Quantiphi, Inc. | Mumbai, India **Oct 2020 – Aug 2021**
- Researched and presented highlights of the three US pandemic stimulus bills to internal stakeholders that informed Quantiphi’s Public Sector business strategy
 - Performed market research on 200 organizations in the US Education industry and came up with an effective go-to market strategy that converted four cold leads
 - Initiated and led the creation of an internal repository to keep track of research advancements in machine learning; this was leveraged by 230 people in the organization including founders
 - Presented solution deck to four leads showcasing how machine learning can be incorporated into leads’ existing processes, converting two
 - Analyzed and reported quarterly revenue figures to internal stakeholders using Looker dashboard
- Freelance Android Developer** | IPLit Solutions, LLP | Thane, India **Feb 2020 – Mar 2020**
- Developed and deployed an Android application for handsfree token printing for use in hospitals and clinics
 - Currently in use in *two hospitals across the city*
- Project Intern** | Fractal Analytics | Mumbai, India **Jun 2019 – Jul 2019**
- Implemented the object classification phase of a project that analyzed consumer behavior at stores for a Fortune 500 FMCG company. Built a model for classifying 50 product SKUs in the product range with 80% accuracy
 - Set up a data augmentation and ingestion pipeline for the classifier. Coded a script for scrapping images of representative products from e-commerce websites to augment data

PUBLISHED RESEARCH

- “Simplification with the Transformer - Its Drawbacks” (International Journal of Computer Sciences and Engineering) [\[pdf\]](#)
- “Abalone Age Prediction Problem: A Review” (International Journal of Computer Applications) [\[pdf\]](#)
- “Kernel Regression from Scratch in Python” (Towards Data Science, Medium.com) [\[link\]](#)
- “What Mainstream AI is Not Doing” (Towards Data Science, Medium.com) [\[link\]](#)
- “10 Points to Make it Big in the Data Industry” (Towards Data Science, Medium.com) [\[link\]](#)
- “Introduction to PySpark via AWS EMR and Hands-on EDA” (Towards AI, Medium.com) [\[link\]](#)

- “Fantasy Premier League x Data Analysis: Being Among the Top 2%” (Towards Data Science, Medium.com) [\[link\]](#)
- “How Gradient Descent Works?” (Towards Data Science, Medium.com) [\[link\]](#)
- “Five Minute Paper Explanations: Part I” (Towards Data Science, Medium.com) [\[link\]](#)
- “Five Minute Paper Explanations: Part II” (Towards Data Science, Medium.com) [\[link\]](#)
- “Five Minute Paper Explanations: Part III” (Towards Data Science, Medium.com) [\[link\]](#)
- “Five Minute Paper Explanations: Part IV” (Towards Data Science, Medium.com) [\[link\]](#)

PROJECTS

- Food AI | Multi-modal Representation Learning** (Python, PyTorch, Hugging Face, seaborn) **May 2022**
- Performed retrieval of food recipes given food images using crossmodal techniques and beat the baseline top-10 recall for recipe retrieval in the original [im2recipe](#) paper by 20 percentage points by improving the feature extraction pipeline
 - Improved recipe retrieval further by projecting learnt features to a shared space using triplet loss trained neural networks and obtained median retrieval rank of 1 and top-10 recall of 82.49% for 1,000 random food images [\[Repo link\]](#)
- Movie Recommendation from Conversational Data | NLP** (Python, PyTorch, surprise, seaborn) **May 2022**
- Built a movie recommendation system leveraging user conversations, critics data and domain adaptation techniques, which is a re-implementation of [this paper](#)
 - Tuned hyperparameters for three CF approaches: KNN, SVD and SVDpp to obtain a 3% improvement in results
 - Experimented with neural MF and obtained comparable results of RMSE=1.232 and MAE=0.9569 [\[Repo link\]](#)
- Logo Detection and Classification | Computer Vision** (Python, PyTorch, seaborn) **May 2022**
- Reproduced results for open set logo detection from the paper [here](#) achieving a 24-percentage point increase in mean average precision compared to the original using YOLOv5
 - Focused on classifying textual logos and obtained a classification accuracy of 22.56% against 47 classes of the Flickr-47 dataset using a logo classification architecture consisting of YOLOv5 and template matching [\[Repo link\]](#)
- Image Colorization using Autoencoders | Convolutional Neural Networks** (Python, Pytorch, streamlit, matplotlib) **Nov 2021**
- Built a 11-layer deep autoencoder neural network using residual connections that colorizes black and white images
 - Trained the network on 10,000 images from FloydHub and deployed online via Streamlit [\[Repo link\]](#)
- New York Taxi Fare Prediction | Big Data** (Python, pandas, matplotlib, PySpark, AWS EMR, AWS EC2) **Oct 2020**
- Performed feature engineering to focus on trips to and from airports and across different boroughs of NYC
 - Predicted taxi fares to a RMSE score of 4.28 by training a Random Forest on the augmented data [\[Repo link\]](#)
- FPL Teammaker | Data Analysis** (Python, NumPy, pandas, streamlit, matplotlib, PuLP) **Sep 2020**
- Developed and deployed an application that performs exploratory data analysis on the English soccer Fantasy Premier League (FPL) game data to suggest an optimal team to be entered into the game
 - 50+ monthly active users. Ranked top 2% in worldwide ranking among 8.2 million players in the year 2020 [\[Repo link\]](#)
- Final Year Project | Natural Language Processing** (Python, Pytorch, Flask, Selenium) **Jun 2019**
- Built a text simplification system that can work on text and simplify it by removing difficult-to-understand words
 - Modeled and trained Transformer models that internalized the semantics of and recognized complex words in input
 - Improved the performance of the application by preceding the transformer architecture with a Complex Word Identification (90.23% accuracy) model that flagged the complex words beforehand [\[Repo link\]](#)
- Abalone Age Prediction | Machine Learning – Regression** (Python, pandas, NumPy, matplotlib, scikit-learn) **Sep 2019**
- Determined the ages of abalones (snails) using classification techniques and leveraging their physical characteristics
 - Improved the accuracy of determining age using regression techniques and obtained a MAE of 0.936
 - Concluded that the dataset is not large enough to get the desired MAE of 0.5 implying correct age prediction [\[Repo link\]](#)
- Alien Shooter | Python Game Development** (Python) **Feb 2019**
- Expanded the ‘Space Invader’ game to include three modes of play: Arcade, Timed and Survival [\[Repo link\]](#)
- Reminder – Todo List | Android Development** (Android, Java) **Jun 2018**
- Developed an Android application that acts as a combination of a reminder app and a notes app
 - Published the app on Google Play Store, and currently has 50+ installs with a rating of 4.6 [\[App link\]](#)

ADDITIONAL ACTIVITIES

- Rehabilitated 10 runaway children as a volunteer for an NGO in Thane, India
- Mentored a team of 12 photographers through covering three college festivals
- Written 4 articles for Kshitij, the KJ Somaiya College of Engineering college magazine
- Visited Vechtdal College, Netherlands as a cultural exchange student